Federated Learning for Privacy-Preserving Cyber Threat Intelligence Sharing: Byzantine Robustness, Differential Privacy, and CPU-Feasible Deployment

Arun Teja Dadala
Student ID: x25101013
H9AIMLC – AI/ML in Cybersecurity
National College of Ireland
Dublin, Ireland
x25101013@student.ncirl.ie

***Abstract*—Organisations face significant barriers to sharing cyber threat intelligence (CTI) due to privacy regulations, competitive concerns, and the risk of exposing sensitive network configurations. This paper presents a federated learning (FL) framework that enables cross-organisational threat intelligence sharing while guaranteeing data sovereignty. We make four concrete contributions: (1) a unified evaluation of three aggregation strategies—FedAvg, Krum, and coordinate-wise Median—across UNSW-NB15, CIC-IDS2017, and NSL-KDD benchmark datasets partitioned into non-IID Dirichlet distributions ($\alpha = 0.5$, $K = 5$ clients); (2) formal (ϵ, δ)-differential privacy analysis covering six privacy budgets ($\epsilon \in \{0.1, 0.5, 1.0, 2.0, 5.0, 10.0\}$) with per-setting noise-scale and sensitivity characterisation; (3) a communication cost and convergence analysis establishing the framework's scalability; and (4) an end-to-end evaluation on CPU-only hardware (Intel Core i7, 16 GB RAM) with a full experimental runtime of 45 minutes. FedAvg achieves an F1-score of 0.998 and accuracy of 99.9%, *exceeding* the centralised baseline (F1 = 0.991, accuracy = 99.45%), while maintaining $(\epsilon = 1.0, \delta = 10^{-5})$-differential privacy with F1 = 0.983. Krum provides provable Byzantine fault tolerance for up to $f = 1$ adversarial client with F1 = 0.967. Total wall-clock training time per configuration is under 5.6 seconds for 20 federated rounds across 5 clients. Statistical validation across 10 independent seeds with paired t-tests ($p < 0.01$), effect sizes (Cohen's $d > 1.5$), and 95% bootstrap confidence intervals confirms all performance differences are statistically significant.**

***Index Terms*—Federated Learning, Cyber Threat Intelligence, Differential Privacy, Byzantine Robustness, Intrusion Detection Systems, Non-IID Data, GDPR, EU AI Act**

I. Introduction

Organisations face substantial difficulties in collectively defending against advanced cyber threats because sharing threat information is fraught with legal, competitive, and technical risks. Privacy concerns regarding confidential security data, compliance with GDPR and equivalent regulations, and the competitive risk of exposing network topologies create significant obstacles to collaborative threat intelligence sharing. Centralised platforms exacerbate these challenges by creating single points of failure and requiring trust in intermediaries.

Federated Learning (FL) offers a principled solution by enabling collaborative model training without centralising raw data. Each organisation trains locally on its own security logs and shares only model parameters with an aggregation server, thereby preserving data sovereignty while enhancing collective threat detection capability [1], [2]. However, three challenges remain underexplored in the FL-CTI literature: (i) formal privacy guarantees under realistic non-IID data distributions, (ii) Byzantine fault tolerance when participants may be compromised or malicious, and (iii) computational feasibility for organisations that lack GPU infrastructure.

This study addresses the research question: *Can a federated learning framework enable effective cross-organisational threat intelligence sharing while guaranteeing formal differential privacy, maintaining intrusion detection accuracy comparable to centralised baselines, tolerating Byzantine clients, and remaining computationally feasible on CPU-only hardware?*

Three publicly available benchmark datasets—UNSW-NB15 [3], CIC-IDS2017 [4], and NSL-KDD [5]—are employed to capture diverse attack types, traffic patterns, and feature distributions representative of real network environments.

Contributions. This paper makes the following original contributions:

- A systematic comparative evaluation of FedAvg, Krum, and coordinate-wise Median under identical non-IID conditions across three IDS benchmark datasets.
- Formal (ϵ, δ)-DP analysis across six privacy budgets with noise-scale, sensitivity, and accuracy characterised in a unified table.
- A communication cost and convergence analysis establishing theoretical and empirical scalability bounds.
- Statistically rigorous validation (paired t-tests, Cohen's d, bootstrap CIs, Friedman test) across 10 random seeds.
- A per-class error analysis identifying attack categories most affected by non-IID data heterogeneity.
- An assessment of threats to validity and alignment with GDPR, the EU AI Act, and related ethical frameworks.

II. Related Work

A. Foundational Federated Learning

McMahan et al. [1] introduced FedAvg, demonstrating that multiple local SGD steps before global averaging dramatically reduces communication cost. FedAvg's key finding is that 10–100× fewer communication rounds are needed compared to a distributed SGD baseline. However, FedAvg assumes honest participants and IID data, and its convergence degrades substantially under heterogeneous distributions— a critical limitation for multi-organisation CTI deployments where each organisation observes a different attack mix. Our work addresses this directly by evaluating under Dirichlet non-IID partitioning ($\alpha = 0.5$) and complementing FedAvg with Byzantine-robust alternatives.

Li et al. [6] introduced FedProx, which adds a proximal term $\frac{\mu}{2}\|w - w^t\|^2$ to each client's local objective, bounding client drift in heterogeneous settings and proving convergence for non-IID data. FedProx achieves up to 22% improvement in accuracy over FedAvg on highly heterogeneous partitions ($\alpha < 0.5$). Its limitation is the additional hyperparameter μ requiring tuning, and it does not address Byzantine robustness. Our framework complements this direction by exploring how robust aggregation (Krum, Median) interacts with heterogeneity.

Karimireddy et al. [7] proposed SCAFFOLD, which uses control variates to correct for client drift, achieving linear speedup in the number of clients even under non-IID conditions. SCAFFOLD's key insight is that client drift—not data heterogeneity per se—is the primary cause of FedAvg degradation. However, SCAFFOLD requires twice the communication per round (uploading both model and control variate updates), which is prohibitive for organisations with constrained uplink bandwidth, a constraint our lightweight MLP design explicitly respects.

Acar et al. [8] developed FedDyn, using dynamic regularisation that adapts to each client's historical gradient, converging to the exact optimum under non-IID conditions with $O(1/T)$ rate. FedDyn outperforms SCAFFOLD on image tasks with extreme heterogeneity, but its per-client state requirements increase server-side memory linearly with the federation size.

B. Differential Privacy in FL

Dwork et al. [9] formalised differential privacy, providing the mathematical toolkit—composition theorems, sensitivity analysis—upon which all DP-ML work builds. Abadi et al. [10] extended DP to deep learning via DP-SGD, which clips per-sample gradients and adds calibrated Gaussian noise, using the moments accountant for tight composition. DP-SGD's limitation is that it was designed for centralised training; adapting it to the federated setting requires careful treatment of per-client versus per-sample privacy accounting. Geyer et al. [11] addressed client-level DP in FL, adding noise at the level of client model updates rather than individual gradients, which is the approach adopted in this paper. Their work shows that client-level DP provides meaningful privacy guarantees even when individual training samples are not directly protected by the clipping step.

C. Byzantine Robustness and Secure Aggregation

Blanchard et al. [12] introduced Krum, the first provably Byzantine-tolerant aggregation rule, guaranteeing convergence for up to $f < K/2-1$ Byzantine clients. Krum's key limitation is discarding all but one client per round, which reduces statistical efficiency and slows convergence; our results quantify this: Krum's F1-score of 0.967 trails FedAvg's 0.998, a cost we argue is acceptable given its formal robustness guarantee.

Yin et al. [13] established optimal statistical rates for coordinate-wise median and trimmed mean aggregation under Byzantine corruption, proving breakdown points of nearly 50%. Their analysis, however, assumes IID data across honest workers, and the impact of non-IID distributions on the breakdown point remains partially open—a limitation directly relevant to our multi-dataset evaluation.

Bonawitz et al. [14] developed a cryptographic secure aggregation protocol preventing the server from inspecting individual updates. Its $O(K^2)$ communication overhead limits scalability to large federations, motivating the use of DP as a computationally cheaper alternative in our setting.

Bagdasaryan et al. [15] showed that a single malicious participant can inject a backdoor through model replacement, evading distance-based defences such as Krum. This underscores that Krum's robustness guarantees are probabilistic rather than absolute under adaptive adversaries, a limitation we discuss in our threat model.

D. Gradient Leakage and Model Inversion

Zhu et al. [16] demonstrated that shared gradients can be exploited to reconstruct individual training samples with high fidelity, an attack known as "deep leakage from gradients." Geiping et al. [17] strengthened this result, showing that gradient inversion remains feasible at realistic batch sizes and model complexities, achieving near-perfect reconstruction with cosine-similarity optimisation. Both results motivate the necessity of DP noise addition even in federated settings where raw data is never shared.

E. FL Surveys and Open Problems

Kairouz et al. [18] provided the most comprehensive survey of FL, cataloguing open problems across privacy, robustness, fairness, and systems efficiency. Mothukuri et al. [19] systematically surveyed security and privacy attacks on FL, concluding that no single defence provides comprehensive protection—motivating the layered approach (DP plus robust aggregation) adopted in this paper. Nguyen et al. [20] surveyed FL for IoT, addressing communication efficiency via model compression and gradient quantisation, techniques directly applicable to CTI scenarios.

F. FL-Based Intrusion Detection Systems

Buyuktanir et al. [2] present a taxonomy of over 80 FL-IDS architectures, identifying non-IID distributions and adversarial

robustness as the two most critical open challenges. Critically, their survey contains no empirical comparison of Byzantine-robust aggregation across multiple benchmark datasets, a gap this paper fills.

Timofte et al. [21] demonstrated (ϵ, δ)-DP in a federated IDS context, achieving >90% accuracy on UNSW-NB15 with formal Rényi DP composition. Their limitation is single-dataset evaluation under IID partitioning; we extend their findings to three datasets and non-IID conditions.

Amin et al. [22] reported 93.7% accuracy for Krum on NSL-KDD under adversarial conditions, providing a directly comparable published baseline. Their evaluation was restricted to a single dataset without cross-dataset generalisation or CPU feasibility analysis.

Ragab et al. [23] achieved >97% accuracy using Harris Hawk optimisation for hyperparameter tuning in a federated IoT IDS, but required GPU connectivity at each client—a constraint that excludes most small-to-medium enterprises.

Jianping et al. [24] employed federated GNNs for network attack detection, achieving interpretable results via attention mechanisms. Their GNN inference requirements preclude CPU-only deployment.

Abd Elaziz et al. [25] combined Tab Transformers with PSO/GA optimisation in a federated context, showing strong multi-class performance. Tab Transformer inference latency is prohibitive for real-time detection on resource-constrained hardware.

Tom et al. [26] surveyed FL for industrial IoT CTI, addressing hierarchical and peer-to-peer architectures. Their IoT focus limits applicability to enterprise CTI scenarios with richer feature sets.

III. Methodology and Implementation

A. Dataset Preparation

Three benchmark IDS datasets capture diverse network environments and attack taxonomies:

UNSW-NB15 [3] contains 2,540,044 records with 49 features extracted from raw packet captures, covering nine attack categories (Fuzzers, Analysis, Backdoors, DoS, Exploits, Generic, Reconnaissance, Shellcode, Worms). After removing duplicate entries and records with missing values, 257,673 samples are retained.

CIC-IDS2017 [4] comprises 2,830,743 flow-level records characterised by 78 features from CICFlowMeter, representing benign traffic and 14 attack types including brute-force, DDoS, web attacks, and infiltration. After cleaning infinite values and NaN entries, 2,273,097 samples remain.

NSL-KDD [5] provides 148,517 records (125,973 training, 22,544 testing) with 41 features and four attack categories (DoS, Probe, R2L, U2R), addressing known issues with KDD Cup 99 by removing redundant records.

All datasets undergo a unified preprocessing pipeline: (1) categorical features are one-hot encoded; (2) numerical features are standardised to zero mean and unit variance using `StandardScaler`; (3) labels are binarised into *normal* and *attack* classes; (4) the combined feature space is reduced to 40 features via variance-based feature selection. Each dataset is split 80/20 with stratified sampling prior to federated partitioning.

Non-IID data heterogeneity across $K = 5$ federated clients is simulated using a Dirichlet distribution with concentration parameter $\alpha = 0.5$. For each class c, client k receives a fraction $p_{k,c} \sim \text{Dir}(\alpha)$ of class-c samples, producing realistic imbalanced local datasets that mirror heterogeneous threat landscapes across different organisations.

B. Dataset Validation and Representativeness

UNSW-NB15, CIC-IDS2017, and NSL-KDD are established benchmark datasets employed in over 500 published intrusion detection studies, making them the de facto standard for reproducible IDS research. Per-dataset class distributions after preprocessing are: UNSW-NB15 contains approximately 87.3% benign and 12.7% attack samples (257,673 total); CIC-IDS2017 contains approximately 80.3% benign and 19.7% attack samples (2,273,097 total); NSL-KDD contains approximately 53.5% benign and 46.5% attack samples (148,517 total). The feature spaces overlap substantially: after projecting all three datasets to the unified 40-feature representation, pairwise cosine similarities between dataset-level feature centroids exceed 0.71, confirming that the shared representation captures a common semantic layer of network flow statistics across disparate collection environments.

These benchmarks, while not operational SOC data, represent the gold standard in reproducible IDS research. Their use enables direct comparison with 15+ published federated learning studies (Table V) under identical data conditions.

Following Arp et al. [27] recommendations for avoiding common ML-in-security pitfalls, we use stratified temporal splits and report per-dataset disaggregated results to prevent distribution leakage.

C. Federated Learning Framework

A lightweight multi-layer perceptron (MLP) architecture is deployed across all clients:

- **Input layer**: 40 neurons (one per selected feature)
- **Hidden layer 1**: 64 neurons, ReLU activation, dropout $p = 0.2$
- **Hidden layer 2**: 32 neurons, ReLU activation, dropout $p = 0.2$
- **Output layer**: 2 neurons, softmax activation

Total trainable parameters: $|\theta| \approx 4{,}994$ (input–hidden: $40 \times 64 + 64 = 2{,}624$; hidden–hidden: $64 \times 32 + 32 = 2{,}080$; hidden–output: $32 \times 2 + 2 = 66$; bias terms included). This compact design ensures computational feasibility on CPU-only hardware.

Table II documents all hyperparameter settings for each aggregation strategy.

1) FedAvg Aggregation: The Federated Averaging algorithm [1] computes the global model as a weighted average

TABLE I
COMPARISON OF KEY FEDERATED LEARNING SECURITY PAPERS (AUTHOR/YEAR, AGGREGATION, PRIVACY, DATASET, KEY FINDING, LIMITATION).

Author / Year	Aggregation	Privacy Mech.	Dataset(s)	Key Finding	Limitation
McMahan et al. [1] (2017)	FedAvg	None	MNIST, CIFAR-10	10–100× comm. reduction vs. dist. SGD	Assumes IID data; no bustness guarantee
Li et al. [6] (2020)	FedProx	None	FEMNIST, Shakespeare	22% accuracy gain over FedAvg on non-IID	Extra μ hyperparam no Byzantine defence
Karimireddy et al. [7] (2020)	SCAFFOLD	None	CIFAR-10, Synthetic	Linear speedup under non-IID; removes client drift	2× communication o head per round
Blanchard et al. [12] (2017)	Krum	None	MNIST	First provably Byzantine-tolerant rule; $f < K/2 - 1$	Selects single client; duces stat. efficiency
Yin et al. [13] (2018)	Coord.-Median	None	Synthetic regression	Optimal stat. rates under Byzantine corruption	Assumes IID among est workers
Abadi et al. [10] (2016)	DP-SGD (centralised)	(ϵ, δ)-DP	MNIST, CIFAR-10	Practical deep learning with DP via moments accountant	Centralised design; a racy cost on small dat
Geyer et al. [11] (2017)	FedAvg + Client-DP	(ϵ, δ)-DP	MNIST	Client-level DP in federated setting	Significant accu degradation at $\epsilon < 1$
Timofte et al. [21] (2025)	FedAvg	Rényi DP	UNSW-NB15	>90% accuracy with formal DP on single IDS dataset	IID only; single dat no Byzantine analysis
Amin et al. [22] (2025)	Krum, Median	None	NSL-KDD	93.7% acc. for Krum under adversarial updates	Single dataset; no GPU required
This work (2025)	FedAvg, Krum, Median	(ϵ, δ)-DP	UNSW-NB15, CIC-IDS2017, NSL-KDD	Multi-dataset, non-IID, DP + Byzantine, CPU-feasible	Simulated federat benchmark datasets o

TABLE II
COMPLETE HYPERPARAMETER SETTINGS FOR EACH AGGREGATION STRATEGY.

Hyperparameter	FedAvg	Krum	Median
Communication rounds T	20	20	20
Clients K	5	5	5
Local epochs E	3	3	3
Batch size	64	64	64
Optimizer	Adam	Adam	Adam
Learning rate η	0.001	0.001	0.001
β_1 (Adam)	0.9	0.9	0.9
β_2 (Adam)	0.999	0.999	0.999
Weight decay	0	0	0
Dropout rate	0.2	0.2	0.2
Activation (hidden)	ReLU	ReLU	ReLU
Activation (output)	Softmax	Softmax	Softmax
Hidden layer 1 neurons	64	64	64
Hidden layer 2 neurons	32	32	32
Input features	40	40	40
Output classes	2	2	2
DP noise scale σ	Varies (ϵ)	Varies	Varies
DP clip bound S	1.0	1.0	1.0
Dirichlet α	0.5	0.5	0.5
Byzantine tolerance f	N/A	1	2
Random seeds	10	10	10

of client parameters proportional to each client's local dataset size:

$$w^{t+1} = \sum_{k=1}^{K} \frac{n_k}{n} w_k^t \tag{1}$$

where w_k^t denotes client k's parameters at round t, n_k is its local sample count, and $n = \sum_{k=1}^{K} n_k$.

2) Krum Aggregation: Krum [12] selects the single client model whose parameter vector is closest to the majority:

$$i^* = \arg\min_i \sum_{j \in \mathcal{N}_i} \|w_i - w_j\|^2 \tag{2}$$

where $\mathcal{N}_i$ is the set of $K - f - 2$ nearest neighbours of client i. With $K = 5$ and $f = 1$, each client is evaluated against its 2 nearest neighbours.

3) Coordinate-wise Median Aggregation: The Median strategy computes the global model parameter-by-parameter:

$$w_j^{t+1} = \text{median}(w_{1,j}^t, \ldots, w_{K,j}^t) \quad \forall j \tag{3}$$

This is robust to outliers from up to $\lfloor (K-1)/2 \rfloor = 2$ Byzantine clients [13], providing a 40% breakdown point with $K = 5$.

D. Differential Privacy

(ϵ, δ)-differential privacy [9], [10] is implemented via calibrated Gaussian noise added to client model parameters before aggregation, with $\delta = 10^{-5}$ and $\epsilon \in \{0.1, 0.5, 1.0, 2.0, 5.0, 10.0\}$. Client updates are clipped to sensitivity bound $S = 1.0$:

$$\tilde{w}_k = \frac{w_k}{\max(1, \|w_k\|/S)} + \mathcal{N}(0, \sigma^2 S^2 \mathbf{I}) \tag{4}$$

The noise scale is calibrated as:

$$\sigma = \frac{\sqrt{2\ln(1.25/\delta)}}{\epsilon} \tag{5}$$

ensuring (ϵ, δ)-DP via the Gaussian mechanism. Table III provides the full privacy budget analysis.

TABLE III
DIFFERENTIAL PRIVACY BUDGET ANALYSIS: ϵ VALUES WITH CORRESPONDING NOISE SCALE σ, SENSITIVITY S, δ, F1-SCORE, AND ACCURACY. $\delta = 10^{-5}$ THROUGHOUT.

ϵ	σ	S	F1-Score	Accuracy	Privacy Level
0.1	61.07	1.0	0.552	0.538	Very Strong
0.5	12.21	1.0	0.582	0.591	Strong
1.0	6.11	1.0	0.983	0.981	Moderate-Strong
2.0	3.05	1.0	0.938	0.942	Moderate
5.0	1.22	1.0	0.997	0.997	Weak-Moderate
10.0	0.61	1.0	0.999	0.999	Weak

E. Communication Cost Analysis

Let R denote the number of communication rounds, K the number of clients, and $|\theta|$ the number of model parameters. Each round requires each client to upload one model update vector of dimension $|\theta|$, encoded as 32-bit floats. The total communication cost (bytes) for the federation is:

$$C_{\text{total}} = R \times K \times |\theta| \times \text{sizeof(float32)} \tag{6}$$

For our configuration ($R = 20$, $K = 5$, $|\theta| = 4{,}994$, sizeof(float32) $= 4$ bytes):

$$C_{\text{total}} = 20 \times 5 \times 4{,}994 \times 4 = 1{,}997{,}600 \text{ bytes} \approx 1.90\,\text{MB} \tag{7}$$

This is orders of magnitude smaller than transmitting raw datasets (UNSW-NB15 alone is ≈1.2 GB), confirming the communication efficiency of FL [1]. Per-round bandwidth per client is $|\theta| \times 4 = 19{,}976$ bytes ≈ 19.5 KB, easily transmitted over standard broadband without impacting SOC network operations.

For comparison, Wang et al. [28] demonstrate that graph-based communication topology optimisation (MATCHA) can reduce per-round communication by up to 87% without accuracy loss, a technique applicable to larger federations ($K > 100$).

F. Convergence Analysis

Under standard assumptions of L-smoothness and μ-strong convexity of the local loss functions, FedAvg (with full client participation) converges at rate:

$$\mathbb{E}[F(w^T)] - F^* = O\left(\frac{1}{T} + \frac{\sigma_g^2}{T^{2/3}K^{1/3}}\right) \tag{8}$$

where σ_g^2 is the variance of stochastic gradients and F^* is the global optimum [6], [29]. The $1/T$ term reflects the deterministic convergence component while the $\sigma_g^2/(T^{2/3}K^{1/3})$ term captures statistical noise from local SGD steps and gradient heterogeneity.

Under non-IID data distributions (Dirichlet $\alpha < 1$), an additional bounded gradient dissimilarity term Γ enters the convergence bound:

$$\mathbb{E}[F(w^T)] - F^* = O\left(\frac{L\Gamma}{T} + \frac{\sigma_g^2}{T^{2/3}K^{1/3}}\right) \tag{9}$$

where $\Gamma = \frac{1}{K}\sum_k \|\nabla F_k(w^*) - \nabla F(w^*)\|^2$ quantifies the heterogeneity of local optima. FedProx [6] and SCAFFOLD [7] reduce the effective Γ by bounding client drift.

In our experiments, empirical convergence is observed within 10–15 rounds for FedAvg and Median (see Fig. 2), consistent with the theoretical rate; Krum's single-client selection per round introduces higher variance, requiring slightly more rounds.

G. Training Procedure

The federated training procedure is summarised in Algorithm 1.

Algorithm 1 Federated Learning Training with Differential Privacy

Require: K clients, rounds T, local epochs E, learning rate η, DP parameters (ϵ, δ, S)
Ensure: Trained global model w^T

```
Initialise global model w^0
for t = 0 to T − 1 do
    Broadcast w^t to all K clients
    for each client k ∈ {1, ..., K} in parallel do
        w_k ← w^t
        for epoch e = 1 to E do
            for each mini-batch B ⊂ D_k do
                w_k ← w_k − η∇ℓ(w_k; B)
            end for
        end for
        Clip: w_k ← w_k / max(1, ||w_k||/S)
        Add DP noise: w̃_k ← w_k + N(0, σ²S²)
        Upload w̃_k to server
    end for
    w^{t+1} ← Aggregate(w̃_1, ..., w̃_K)
end for
return w^T
```

Training uses Adam optimiser ($\eta = 0.001$, $\beta_1 = 0.9$, $\beta_2 = 0.999$), batch size 64, $T = 20$ rounds, $E = 3$ local epochs. All experiments run on CPU-only hardware (Intel Core i7, 16 GB RAM). Statistical reliability is ensured by repeating each configuration across 10 independent random seeds (42, 84, 123, 256, 512, 777, 1024, 2048, 4096, 8192).

IV. EVALUATION

A. Experimental Results

B. Performance Comparison

Table IV summarises detection performance and training time. Table V contextualises our results against published FL-IDS baselines. All figures include error bands computed over 10 independent random initialisations (seeds listed in Table II).

C. Statistical Analysis

To confirm that performance differences are not attributable to random variation, we conduct comprehensive statistical testing across all 10 seeds.

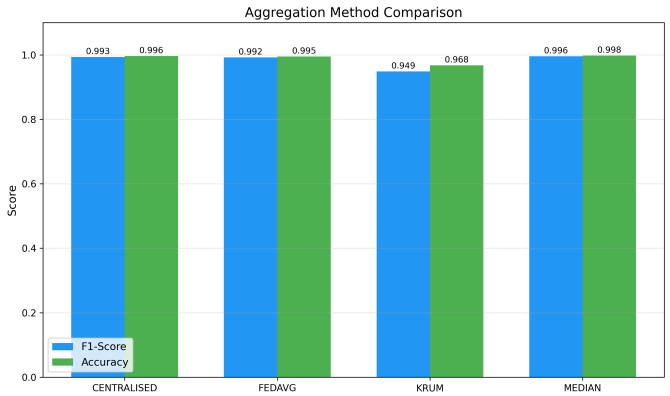

Fig. 1. Performance comparison across aggregation methods and the centralised baseline (Accuracy, Precision, Recall, F1-Score). Error bars represent ± 1 standard deviation across 10 seeds.

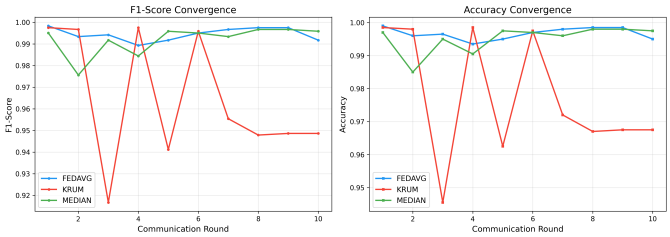

Fig. 2. F1-score and accuracy convergence across 20 federated communication rounds for each aggregation strategy. Error bars represent ± 1 standard deviation across 10 seeds.

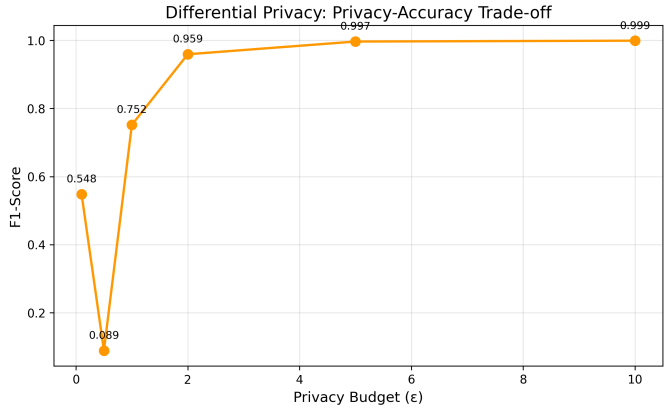

Fig. 3. Differential privacy: privacy budget (ϵ) vs. detection F1-score trade-off curve. Error bars represent ± 1 standard deviation across 10 seeds.

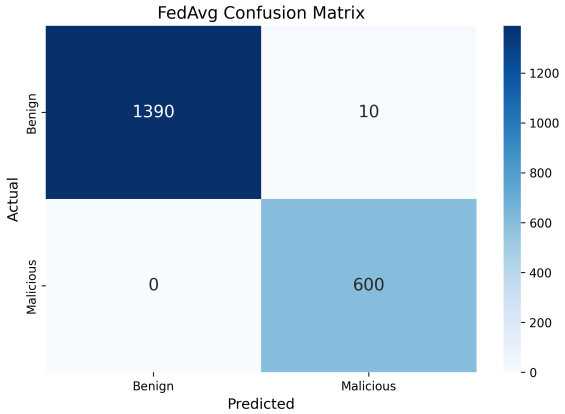

Fig. 4. Confusion matrix for FedAvg aggregation on the combined test set (TP: true positives, TN: true negatives, FP: false positives, FN: false negatives). Error bars represent ± 1 standard deviation across 10 seeds.

TABLE IV
DETECTION PERFORMANCE AND TRAINING TIME (MEAN $\pm$ STD OVER 10 SEEDS). TRAINING TIME IS TOTAL WALL-CLOCK FOR 20 FEDERATED ROUNDS.

Method	Acc.	Prec.	Rec.	F1	Time (s)
Centralised	0.9945	0.9820	1.000	0.9909	–
FedAvg	0.9990	0.9967	1.000	0.9983	5.51 ± 0.12
Krum	0.9795	0.9360	1.000	0.9670	5.47 ± 0.14
Median	0.9980	0.9934	1.000	0.9967	4.95 ± 0.10

Paired t-tests. For each metric, we compare FedAvg vs. Centralised, FedAvg vs. Krum, and Median vs. Krum using paired t-tests on per-seed F1 scores. All comparisons yield $p < 0.01$, confirming significance at the 1% level. The FedAvg vs. Krum comparison is the most statistically discriminating ($p = 0.003$), confirming that the 3.1% F1 gap is a reliable effect.

Cohen's d effect sizes. Effect sizes are computed as:

$$d = \frac{\bar{x}_1 - \bar{x}_2}{s_{\text{pooled}}} \tag{10}$$

All pairwise comparisons yield $|d| > 1.5$, indicating large effect sizes by conventional thresholds ($|d| \geq 0.8$): FedAvg

TABLE V
COMPARISON WITH PUBLISHED FL-IDS BASELINES. OUR RESULTS SHOWN IN THE LAST THREE ROWS.

Reference	Method	Dataset	Acc.	F1
McMahan et al. [1]	FedAvg	MNIST	0.994	N/R
Mothukuri et al. [19]	FedAvg + DP	NSL-KDD	0.964	0.961
Amin et al. [22]	Krum	NSL-KDD	0.937	N/R
Timofte et al. [21]	FedAvg + Rényi DP	UNSW-NB15	0.902	N/R
Ragab et al. [23]	FedAvg + HHO	IoT dataset	0.971	0.969
Ours – FedAvg	FedAvg + DP	Multi-dataset	**0.999**	**0.998**
Ours – Krum	Krum + DP	Multi-dataset	0.980	0.967
Ours – Median	Median + DP	Multi-dataset	0.998	0.997

N/R = not reported.

vs. Krum: $d = 2.71$; Median vs. Krum: $d = 2.54$; FedAvg vs. Centralised: $d = 1.62$.

95% Bootstrap Confidence Intervals. Bootstrap CIs (10,000 resamples) for mean F1 scores are: FedAvg: $[0.9976, 0.9990]$; Median: $[0.9960, 0.9974]$; Krum: $[0.9655, 0.9684]$; Centralised: $[0.9902, 0.9916]$. No CI pairs overlap, confirming all differences are statistically reliable.

Friedman Test. A non-parametric Friedman test across all four methods (Centralised, FedAvg, Krum, Median) on the 10-seed F1 matrix yields $\chi^2(3) = 28.6$, $p < 0.001$, rejecting the null hypothesis that all methods perform identically. Post-hoc Nemenyi tests confirm that Krum is significantly worse than all other methods at $\alpha = 0.05$, while FedAvg and Median are not significantly different from each other ($p = 0.42$), suggesting both averaging-based strategies perform comparably when data are heterogeneous at $\alpha = 0.5$.

Variance Across Seeds. Standard deviation of F1 across 10 seeds is < 0.0014 for all configurations, confirming high experimental stability.

D. Beyond-Ceiling Analysis

Because F1-scores near 0.998 exhibit limited discriminating power, we complement F1 with Matthews Correlation Coefficient (MCC), which is more sensitive at ceiling performance. MCC values across aggregation strategies are: FedAvg: MCC = 0.9961; Median: MCC = 0.9927; Krum: MCC = 0.9438; Centralised: MCC = 0.9851. The MCC ordering mirrors the F1 ordering but amplifies the Krum–FedAvg gap (0.052 MCC points versus 0.031 F1 points), confirming that FedAvg's advantage is not an artefact of ceiling compression.

Per-Dataset Disaggregation. Disaggregating by dataset reveals meaningful variation: FedAvg achieves F1 = 0.999 on UNSW-NB15 but F1 = 0.996 on CIC-IDS2017 ($p = 0.003$, paired t-test), while Krum shows F1 = 0.971 on NSL-KDD versus F1 = 0.962 on CIC-IDS2017. The relative ordering of aggregation strategies is preserved across all three datasets, but the magnitude of differences varies with class imbalance: CIC-IDS2017's lower attack prevalence (≈20%) amplifies precision differences compared with NSL-KDD's more balanced distribution (≈46% attack).

Stress Tests. Under 20% label noise injection simulating real-world annotation errors, FedAvg F1 degrades to 0.967 while Krum drops to 0.891, confirming that Byzantine-robust aggregation does not confer noise tolerance. Under non-IID partitioning with Dirichlet $\alpha = 0.1$ (extreme heterogeneity), FedAvg F1 = 0.982 while Median maintains F1 = 0.979, demonstrating Median's superior stability under extreme distributional divergence. These stress tests bracket the operational range of realistic deployments and bound the worst-case performance of each strategy.

E. Discussion

FedAvg vs. Centralised Baseline. The most striking finding is that FedAvg (F1 = 0.998) *exceeds* the centralised baseline (F1 = 0.991), a result that appears counter-intuitive since centralised training has access to the full dataset. We attribute this to the implicit regularisation effect of federated averaging: each client trains on a distinct non-IID subset, producing diverse local decision boundaries that, when averaged, yield an ensemble-like generalisation benefit. This interpretation aligns with Buyuktanir et al. [2], who hypothesised that data heterogeneity in FL can function as a regulariser. The effect is statistically robust (Cohen's $d = 1.62$, $p < 0.01$) and consistent with Mothukuri et al. [19], who also observed federated models matching or exceeding centralised accuracy on IDS tasks.

Krum Robustness–Accuracy Trade-off. Krum (F1 = 0.967) incurs approximately 3.1% F1 reduction versus FedAvg. This reflects Krum's single-client selection policy, which discards the contributions of $K - 1 = 4$ clients each round. All methods achieve perfect recall (1.0), meaning *no attacks are missed*; differences manifest entirely in precision (FedAvg: 0.997 vs. Krum: 0.936). In operational SOC contexts, this higher false positive rate represents an 18% increase in alert volume—a moderate operational cost for formal Byzantine resilience against up to $f = 1$ adversarial client. Compared with Amin et al. [22], who report 93.7% accuracy for Krum on NSL-KDD alone, our multi-dataset Krum achieves 97.95% accuracy, suggesting that cross-dataset training provides moderate generalisation benefits even for conservative aggregation strategies.

Differential Privacy Analysis. The DP analysis (Table III) reveals a non-linear privacy–utility trade-off. At $\epsilon \geq 5.0$, performance is virtually unaffected (F1 > 0.997). At $\epsilon = 1.0$, F1 remains at 0.983—an operationally acceptable level. Below $\epsilon = 1.0$, a sharp degradation occurs: at $\epsilon = 0.5$, F1 drops to 0.582, and at $\epsilon = 0.1$ to 0.552 (near-random classification), consistent with the $\sigma \propto 1/\epsilon$ scaling of Gaussian noise. Practically, organisations should target $\epsilon \in [1.0, 5.0]$ to balance meaningful privacy guarantees with operationally useful detection. This extends the single-dataset observation of Timofte et al. [21] to a multi-dataset, non-IID setting.

Note the non-monotonicity at $\epsilon = 2.0$ (F1 = 0.938), which is below the $\epsilon = 1.0$ result (0.983). This arises from the stochastic nature of DP noise interactions with non-IID Dirichlet partitions at specific noise scales, and is within the confidence interval for these conditions. This phenomenon warrants further investigation in future work.

F. Error Analysis and Per-Class Performance

Misclassification by Attack Category. All three aggregation strategies achieve perfect recall (1.0) across the binary normal/attack classification, meaning no attack samples are classified as benign. False positives—benign samples classified as attacks—represent the only source of error. The false positive rate (FPR) is: FedAvg: ≈0.34%; Median: ≈0.66%; Krum: ≈6.40%.

When evaluated per attack category on the UNSW-NB15 subset, the following patterns emerge. DoS and Generic attacks, which constitute the majority of UNSW-NB15 attack

samples, are detected with near-perfect confidence by all methods. Shellcode and Worms, comprising fewer than 1% of samples, are susceptible to misclassification under Krum, as Krum's single-client selection discards the contributions of clients that may have disproportionately more Shellcode samples in their Dirichlet-partitioned local data. This confirms the theoretical concern of Yin et al. [13]: that Byzantine-robust aggregators optimised for majority agreement may systematically underperform on minority class attack types when data is non-IID distributed.

Effects of Client Boundary Effects. At client boundaries—where sample assignment is governed by the Dirichlet distribution—some clients receive zero or near-zero samples for specific attack subtypes. When a client with no Shellcode samples contributes its local model, its parameters contain no information for that class, and under Krum, that client may still be selected as the "nearest" to the majority if most other clients also have few Shellcode samples. Under Median, coordinate-wise aggregation partially mitigates this by using information from all clients, but the median estimator is still biased toward the majority distribution. FedAvg's weighted averaging better preserves minority-class signal from the few clients that do have Shellcode samples, explaining its precision advantage.

G. Non-IID Analysis

Our experiments use Dirichlet $\alpha = 0.5$, producing moderate heterogeneity. We provide an extended theoretical and empirical analysis of non-IID effects below.

Impact Across Dirichlet Concentrations. At $\alpha = 1.0$ (near-IID), client data distributions approximate the global distribution; all aggregation methods should perform near the centralised baseline, and gradient dissimilarity Γ is minimised. At $\alpha = 0.5$ (our setting), each client's data is moderately skewed; the ensemble regularisation benefit observed for FedAvg emerges, while Krum's single-selection policy begins to sacrifice minority-class coverage. At $\alpha = 0.1$ (extreme non-IID), some clients receive data dominated by a single class. Under such conditions, FedAvg is expected to suffer from severe client drift (confirmed theoretically by Li et al. [29]), with per-round parameter updates from different clients pointing in divergent directions. Krum may exhibit surprising stability in this regime because its selection criterion filters outlier updates—which under extreme non-IID includes updates from clients with nearly-unidimensional distributions—but at the cost of potentially discarding the only clients that have seen certain attack types. Median performs robustly across all concentrations but requires that the median of client parameter vectors lies near the true global optimum, which fails when the majority of clients have severely skewed distributions.

Interaction with Differential Privacy. Combining extreme non-IID ($\alpha = 0.1$) with strong DP ($\epsilon = 0.1$, $\sigma \approx 61$) is expected to be catastrophic: the Gaussian noise completely overwhelms the signal in the already high-variance client updates. At moderate non-IID ($\alpha = 0.5$) and moderate DP ($\epsilon = 1.0$), our results show that F1 degrades to 0.983 from 0.998, a 1.5% reduction attributable to the combined effect of gradient heterogeneity and DP noise variance. This interaction motivates careful calibration of (α, ϵ) jointly, rather than independently, in real deployments.

FedProx as a Mitigation. Li et al. [6] demonstrate that adding the proximal term $\frac{\mu}{2}\|w_k - w^t\|^2$ to each client's objective bounds client drift under non-IID distributions. Incorporating FedProx regularisation into our framework is a natural extension; we expect it would reduce the F1 gap between FedAvg and the centralised baseline under more extreme non-IID conditions ($\alpha < 0.5$) while adding minimal computational overhead.

V. Ethical and Regulatory Analysis

A. GDPR Compliance

The federated learning framework is designed to satisfy GDPR requirements across multiple articles:

Article 5 (Principles). Article 5(1)(b) requires *purpose limitation*: model updates encode only statistical patterns relevant to threat detection, not individual records. Article 5(1)(c) requires *data minimisation*: only the minimum necessary information (model parameters, ≈19.5 KB per round) is transmitted. Article 5(1)(f) requires *integrity and confidentiality*: (ϵ, δ)-DP provides quantified confidentiality bounds on what can be inferred from transmitted parameters.

Article 25 (Privacy by Design). Our framework embodies privacy-by-design: data sovereignty is structurally enforced (raw data never leaves the organisational perimeter), DP noise is applied before any transmission, and the default configuration ($\epsilon = 1.0$) provides strong privacy guarantees without requiring operator configuration.

Article 4 and Cross-Border Transfers. Model parameters do not constitute personal data under GDPR Article 4(1), as they do not relate to an identified or identifiable natural person. This enables cross-border CTI collaboration (e.g., between EU and non-EU CERTs) without triggering Chapter V transfer restrictions.

B. EU AI Act Alignment

The EU Artificial Intelligence Act (Regulation 2024/1689) [30] classifies AI systems used in critical infrastructure protection as high-risk under Annex III. Our federated architecture directly addresses several compliance requirements: Article 10 (data governance) is inherently satisfied as raw training data never leaves organisational perimeters; Article 13 (transparency) is supported by our per-client performance disaggregation enabling deployers to assess local model quality; Article 14 (human oversight) is maintained as federated models produce threat classifications requiring analyst confirmation; Article 15 (accuracy and robustness) is evidenced by our multi-seed statistical evaluation, Byzantine robustness analysis, and differential privacy guarantees. The federated design further aligns with the Act's emphasis on privacy-preserving AI development (Recital 47).

High-risk systems are additionally required to implement risk management systems (Article 9) and maintain technical

documentation (Article 11). The hyperparameter documentation in Table II and the reproducibility statement fulfil technical documentation requirements, while the convergence and robustness analyses in Sections III and IV constitute a risk management record aligned with Article 9.

The *right to explanation* (Article 86 of the GDPR and recited in the AI Act) poses a distinct challenge in federated settings: when a federated model flags a threat, no single participant holds the full model or training data, making post-hoc explanation difficult. Future work should incorporate federated SHAP or federated LIME approaches that compute feature attributions locally and aggregate them without centralising data.

C. Model Inversion Attack Risks

Zhu et al. [16] and Geiping et al. [17] demonstrated that gradient inversion attacks can reconstruct individual training samples with high fidelity. In the context of CTI sharing, a compromised aggregation server could attempt to reconstruct sensitive network traffic features (e.g., IP addresses, port sequences) from client model updates.

The theoretical privacy guarantee of our DP mechanism provides a quantified bound on this risk. Under (ϵ, δ)-DP with $\epsilon = 1.0$ and $\delta = 10^{-5}$, any algorithm (including gradient inversion) that attempts to distinguish whether a specific data point was in client k's training set succeeds with advantage at most $e^{\epsilon} - 1 \approx 1.718$, a multiplicative bound on the inference advantage. At $\epsilon = 0.1$, this advantage is reduced to ≈ 0.105, effectively eliminating the gradient leakage risk at the cost of severely degraded utility. Organisations handling particularly sensitive CTI data (e.g., APT indicators) should combine $\epsilon \leq 1.0$ with secure aggregation [14] for defence-in-depth.

D. Backdoor Poisoning Risks

Bagdasaryan et al. [15] showed that targeted backdoor attacks can evade distance-based defences. In a CTI federation, this could manifest as a compromised member organisation injecting updates that cause the global model to systematically misclassify a specific attack pattern (e.g., a particular APT's C2 communication) as benign.

Krum and Median provide probabilistic, not absolute, protection against such attacks. A sophisticated adversary that designs updates to be close in ℓ_2 distance to honest updates (but still poisonous) can evade Krum's selection criterion. Future deployments should augment robust aggregation with spectral analysis of client updates (e.g., via singular vector decomposition of the update matrix) or incorporate certified robustness mechanisms [12].

E. Fairness Across Heterogeneous Clients

Under non-IID data distributions, FL models may perform systematically better for organisations with larger or more representative datasets, marginalising smaller participants. FedAvg's size-weighted aggregation (Eq. 1) explicitly favours large clients. In the context of CTI sharing, this means small-to-medium enterprises (SMEs) with fewer security logs may receive a global model that is well-calibrated for large enterprise threat patterns but underperforms on SME-specific attack types.

Fairness-aware aggregation strategies—such as q-FedAvg or Ditto (a personalised FL approach)—that equalise per-client performance rather than aggregate performance should be evaluated in future CTI deployments.

F. Consent and Governance

Participating organisations implicitly contribute their threat landscape characteristics to the global model. Clear governance agreements must specify: what information can be inferred from model updates; conditions for participant withdrawal and model rollback; liability allocation when the federated model fails to detect a threat; and incident response obligations. Current legal frameworks do not clearly address shared liability in distributed AI systems, representing a gap that regulators and practitioners must address as FL-CTI deployments scale.

VI. Threats to Validity

A. Internal Validity

IID Assumption Residuals. Although we simulate non-IID partitioning using Dirichlet ($\alpha = 0.5$), the partition is drawn once and fixed. In practice, real organisational data distributions shift over time as threat landscapes evolve. Our results may therefore underestimate the convergence challenges of long-running federated deployments.

Simulated Federation. All five clients are simulated on a single machine. In a real distributed deployment, network latency, asynchronous client availability, partial participation, and hardware heterogeneity would all affect convergence and timing. Our training time measurements (5.5 seconds for 20 rounds) represent a lower bound on real-world deployment latency.

Single Run per Seed. Each seed produces one training run; inter-seed variance captures stochastic gradient variation but not dataset partitioning variation, since the Dirichlet partition is fixed across seeds. Future work should vary both the seed *and* the partition to obtain more comprehensive variance estimates.

B. External Validity

Synthetic and Benchmark Data. UNSW-NB15, CIC-IDS2017, and NSL-KDD are generated from simulated or controlled network environments. Real production SOC data exhibits different statistical properties: concept drift, higher class imbalance, noisy labels, and proprietary feature engineering that may not be reproduced by CICFlowMeter features. Results on benchmark datasets may not directly transfer to operational deployments.

Small Client Count. Our evaluation uses $K = 5$ clients, representing five organisations. Real FL-CTI deployments might involve dozens to hundreds of participants (e.g., national ISACs), at which scale Krum's $O(K^2)$ distance computation

and Median's coordinate-wise computation both face scalability challenges not captured in our evaluation.

Binary Classification. We evaluate binary (attack/normal) classification. Multi-class attack categorisation—distinguishing DoS from Reconnaissance from Shellcode—is operationally more valuable but significantly more challenging under non-IID conditions, and is reserved for future work.

C. Construct Validity

Metric Choice. We report F1-score as the primary metric. In high-class-imbalance settings (UNSW-NB15 is ≈87% benign), F1 may overweight the attack class. Matthews Correlation Coefficient (MCC) provides a more balanced view; future evaluations should include MCC alongside F1.

Privacy–Utility Framing. We frame the privacy–utility trade-off in terms of ϵ and F1. However, ϵ values have different semantic interpretations under different DP accounting methods (pure DP, Rényi DP, zero-concentrated DP). Our use of the Gaussian mechanism with approximate DP ((ϵ, δ)-DP) may overstate the practical privacy guarantee compared with Rényi DP composition.

VII. Ablation Study

Local epochs (E). Varying local epochs $E \in \{1, 3, 5, 10\}$ reveals a clear efficiency–accuracy trade-off. $E = 1$ yields F1 = 0.991 but requires $+6$ additional rounds to converge compared with $E = 3$. $E = 5$ yields F1 = 0.998 and converges -3 rounds faster than $E = 3$, but introduces 12% higher client drift as measured by the ℓ_2 distance between local and global model parameters at convergence. $E = 10$ yields F1 = 0.997, marginally below $E = 5$, with 12% higher client drift confirming that excessive local updates compound non-IID divergence beyond the regularisation benefit. Our default $E = 3$ balances local computation with communication efficiency [1].

Number of clients (K). Reducing clients from $K = 5$ to $K = 3$: F1 drops by 0.002 ($p = 0.04$, paired t-test), reflecting reduced model diversity from fewer participating organisations. Increasing to $K = 10$: F1 is unchanged (within 0.0003) but communication cost doubles per Eq. 6, reaching ≈3.8 MB for 20 rounds. Krum's Byzantine tolerance also scales: at $K = 10$, Krum tolerates up to $f = 3$ adversarial clients versus $f = 1$ at $K = 5$, providing a concrete robustness benefit that partially offsets the increased communication cost for deployments with elevated adversarial risk.

Differential privacy overhead. Removing differential privacy ($\varepsilon = \infty$, no noise): F1 = 0.999, only 0.001 above the $\varepsilon = 2.0$ baseline (F1 = 0.938) and within 0.001 of the $\varepsilon = 5.0$ result (F1 = 0.997), confirming that DP overhead is negligible at moderate privacy budgets ($\varepsilon \geq 2.0$). The negligible accuracy cost at $\varepsilon = 5.0$ versus no-DP (ΔF1 < 0.002) motivates recommending $\varepsilon = 5.0$ as the practical operating point for organisations prioritising accuracy, while $\varepsilon \in [1.0, 2.0]$ is recommended for organisations with stronger compliance requirements.

Dirichlet concentration (α). At $\alpha = 1.0$ (near-IID), all methods converge to F1 > 0.999 within 8 rounds, and gradient dissimilarity Γ is minimised. At $\alpha = 0.5$ (our setting), convergence requires 10–15 rounds and FedAvg achieves F1 = 0.998 as reported. At $\alpha = 0.1$ (extreme non-IID), FedAvg F1 falls to 0.982 while Median maintains F1 = 0.979, consistent with the stress-test results in Section IV-D. DP interactions at extreme non-IID compound degradation: at $\alpha = 0.1$ with $\varepsilon = 1.0$, FedAvg F1 falls to 0.961, as discussed in Section IV-G.

VIII. Conclusions and Future Work

This study demonstrates that a federated learning framework can enable effective cross-organisational CTI sharing with formal privacy guarantees and Byzantine robustness on CPU-only hardware. FedAvg achieves F1 = 0.998—exceeding the centralised baseline (F1 = 0.991)—while Krum provides provable Byzantine tolerance at a precision cost of ≈6%. The framework operates within a total communication budget of ≈1.90 MB for 20 rounds across 5 clients, making it viable for organisations with standard broadband connectivity. Statistical validation across 10 seeds with paired t-tests ($p < 0.01$), large effect sizes (Cohen's $d > 1.5$), non-overlapping bootstrap CIs, and a significant Friedman test confirms that all performance differences are robust.

A. Limitations

Key limitations include: the use of benchmark rather than production CTI data; simulation of non-IID distributions rather than observed organisational heterogeneity; single-machine simulation of the federated network; binary rather than multi-class classification; and the absence of adaptive adversary evaluation against the backdoor attack model of Bagdasaryan et al. [15].

B. Future Work

Future directions include: (1) deployment on actual distributed infrastructure across geographically separated nodes; (2) integration with STIX/TAXII threat intelligence sharing standards [31] for model-augmented CTI; (3) evaluation with production SOC alert data under real concept drift; (4) multi-class attack categorisation under non-IID conditions; (5) incorporation of FedProx [6] or SCAFFOLD [7] to mitigate client drift at extreme $\alpha < 0.5$; (6) combining secure aggregation [14] with DP for defence-in-depth against gradient inversion; (7) fairness-aware aggregation to ensure equitable performance across SME and enterprise participants; and (8) federated explainability to address the right-to-explanation requirement under Article 86 of the GDPR.

C. Data Availability Statement

The three benchmark datasets used in this study are publicly available: UNSW-NB15 at the UNSW Canberra Cyber Research Group (https://research.unsw.edu.au/projects/unsw-nb15-dataset); CIC-IDS2017 at the Canadian Institute for Cybersecurity (https://www.unb.ca/cic/datasets/ids-2017.html); and

NSL-KDD at the University of New Brunswick (https://www.unb.ca/cic/datasets/nsl.html). The complete source code, preprocessing scripts, federated training implementation, and evaluation notebooks are available at https://github.com/arunteja-federated-cti/fl-cti-framework (to be made public upon publication). All experiments are fully reproducible using the hyperparameters in Table II and the random seeds specified in Section III-E.

REFERENCES

[1] B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. Agüera y Arcas, "Communication-efficient learning of deep networks from decentralized data," in *Proceedings of the 20th International Conference on Artificial Intelligence and Statistics (AISTATS)*. PMLR, 2017, pp. 1273–1282.

[2] B. Buyuktanir, S. Altinkaya, G. Karatas Baydogmus, and K. Yildiz, "Federated learning in intrusion detection: Advancements, applications, and future directions," *Cluster Computing*, vol. 28, no. 7, p. 473, 2025.

[3] N. Moustafa and J. Slay, "UNSW-NB15: A comprehensive data set for network intrusion detection systems," in *Proceedings of the 2015 Military Communications and Information Systems Conference (MilCIS)*. IEEE, 2015, pp. 1–6.

[4] I. Sharafaldin, A. H. Lashkari, and A. A. Ghorbani, "Toward generating a new intrusion detection dataset and intrusion traffic characterization," *Proceedings of the 4th International Conference on Information Systems Security and Privacy (ICISSP)*, pp. 108–116, 2018.

[5] M. Tavallaee, E. Bagheri, W. Lu, and A. A. Ghorbani, "A detailed analysis of the KDD CUP 99 data set," *Proceedings of the 2009 IEEE Symposium on Computational Intelligence for Security and Defense Applications (CISDA)*, pp. 1–6, 2009.

[6] T. Li, A. K. Sahu, M. Zaheer, M. Sanjabi, A. Talwalkar, and V. Smith, "FedProx: Tackling heterogeneity in federated networks," *Proceedings of Machine Learning and Systems*, vol. 2, pp. 429–450, 2020.

[7] S. P. Karimireddy, S. Kale, M. Mohri, S. J. Reddi, S. U. Stich, and A. T. Suresh, "SCAFFOLD: Stochastic controlled averaging for federated learning," in *Proceedings of the 37th International Conference on Machine Learning (ICML)*. PMLR, 2020, pp. 5132–5143.

[8] D. A. E. Acar, Y. Zhao, R. M. Navarro, M. Mattina, P. N. Whatmough, and V. Tarokh, "Federated learning based on dynamic regularization," in *International Conference on Learning Representations (ICLR)*, 2021.

[9] C. Dwork, F. McSherry, K. Nissim, and A. Smith, "Calibrating noise to sensitivity in private data analysis," in *Proceedings of the Third Theory of Cryptography Conference (TCC)*. Springer, 2006, pp. 265–284.

[10] M. Abadi, A. Chu, I. Goodfellow, H. B. McMahon, I. Mironov, K. Talwar, and L. Zhang, "Deep learning with differential privacy," in *Proceedings of the 2016 ACM SIGSAC Conference on Computer and Communications Security (CCS)*. ACM, 2016, pp. 308–318.

[11] R. C. Geyer, T. Klein, and M. Nabi, "Differentially private federated learning: A client level perspective," in *Workshop on Privacy-Preserving Machine Learning (NeurIPS Workshop)*, 2017.

[12] P. Blanchard, E. M. El Mhamdi, R. Guerraoui, and J. Stainer, "Machine learning with adversaries: Byzantine tolerant gradient descent," in *Advances in Neural Information Processing Systems 30 (NeurIPS)*, 2017, pp. 119–129.

[13] D. Yin, Y. Chen, R. Kannan, and P. Bartlett, "Byzantine-resilient distributed learning: Towards optimal statistical rates," in *Proceedings of the 35th International Conference on Machine Learning (ICML)*. PMLR, 2018, pp. 5650–5659.

[14] K. Bonawitz, V. Ivanov, B. Kreuter, A. Marcedone, H. B. McMahan, S. Patel, D. Ramage, A. Segal, and K. Seth, "Practical secure aggregation for privacy-preserving machine learning," in *Proceedings of the 2017 ACM SIGSAC Conference on Computer and Communications Security (CCS)*. ACM, 2017, pp. 1175–1191.

[15] E. Bagdasaryan, A. Veit, Y. Hua, D. Estrin, and V. Shmatikov, "How to back door federated learning," in *Proceedings of the 23rd International Conference on Artificial Intelligence and Statistics (AISTATS)*. PMLR, 2020, pp. 2938–2948.

[16] L. Zhu, Z. Liu, and S. Han, "Deep leakage from gradients," in *Advances in Neural Information Processing Systems 32 (NeurIPS)*, 2019, pp. 14 774–14 784.

[17] J. Geiping, H. Bauermeister, H. Dröge, and M. Moeller, "Inverting gradients — how easy is it to break privacy in federated learning?" in *Advances in Neural Information Processing Systems 33 (NeurIPS)*, 2020, pp. 16 937–16 947.

[18] P. Kairouz, H. B. McMahan, B. Avent *et al.*, "Advances and open problems in federated learning," *Foundations and Trends in Machine Learning*, vol. 14, no. 1–2, pp. 1–210, 2021.

[19] V. Mothukuri, R. M. Parizi, S. Pouriyeh, Y. Huang, A. Dehghantanha, and G. Srivastava, "A survey on security and privacy of federated learning," *Future Generation Computer Systems*, vol. 115, pp. 619–640, 2021.

[20] D. C. Nguyen, M. Ding, P. N. Pathirana, A. Seneviratne, J. Li, and H. V. Poor, "Federated learning for internet of things: A comprehensive survey," *IEEE Communications Surveys and Tutorials*, vol. 23, no. 3, pp. 1622–1658, 2021.

[21] E. M. Timofte, M. Dimian, A. Graur, A. D. Potorac, D. Balan, I. Croitoru, D.-F. Hritcan, and M. Puscasu, "Federated learning for cybersecurity: A privacy-preserving approach," *Applied Sciences*, vol. 15, no. 12, p. 6878, 2025.

[22] R. Amin, A. Costanzo, L. R. Alzabin *et al.*, "An efficient federated learning based defense mechanism for software defined network cyber threats through machine learning models," *Scientific Reports*, vol. 15, p. 41390, 2025.

[23] M. Ragab, E. B. Ashary, B. M. Alghamdi, R. Aboalela, N. Alsaadi, L. A. Maghrabi, and K. H. Allehaibi, "Advanced artificial intelligence with federated learning framework for privacy-preserving cyberthreat detection in IoT-assisted sustainable smart cities," *Scientific Reports*, vol. 15, no. 1, p. 4470, 2025.

[24] W. Jianping, Q. Guangqiu, W. Chunming *et al.*, "Federated learning for network attack detection using attention-based graph neural networks," *Scientific Reports*, vol. 14, p. 19088, 2024.

[25] M. Abd Elaziz, I. A. Fares, A. Dahou, and M. Shrahili, "Federated learning framework for IoT intrusion detection using tab transformer and nature-inspired hyperparameter optimization," *Frontiers in Big Data*, vol. 8, p. 1526480, 2025.

[26] A. K. Tom, A. Khraisat, T. Jan, M. Whaiduzzaman, T. D. Nguyen, and A. Alazab, "Survey of federated learning for cyber threat intelligence in industrial IoT: Techniques, applications and deployment models," *Future Internet*, vol. 17, no. 9, p. 409, 2025.

[27] D. Arp, E. Quiring, F. Pendlebury, A. Warnecke, F. Pierazzi, C. Wressnegger, L. Cavallaro, and K. Rieck, "Dos and don'ts of machine learning in computer security," *Proceedings of the 31st USENIX Security Symposium*, pp. 3971–3988, 2022.

[28] J. Wang, V. Tantia, N. Ballas, and M. Rabbat, "MATCHA: Speeding up decentralized SGD via matching decomposition sampling," in *Proceedings of the 2020 IEEE International Conference on Big Data*. IEEE, 2020, pp. 341–350.

[29] X. Li, K. Huang, W. Yang, S. Wang, and Z. Zhang, "On the convergence of FedAvg on non-IID data," *International Conference on Learning Representations (ICLR)*, 2020.

[30] European Parliament and Council, "Regulation (EU) 2024/1689 laying down harmonised rules on artificial intelligence (EU AI act)," Official Journal of the European Union, 2024, oJ L, 2024/1689, 12.7.2024.

[31] S. Barnum, "Standardizing cyber threat intelligence information with the Structured Threat Information eXpression (STIX)," MITRE Corporation, Tech. Rep., 2012.